

001191074

April 26, 2022

```
[1]: import matplotlib.pyplot as plt
from PIL import Image
import seaborn as sns
import numpy as np
import pandas as pd
import os
from tensorflow.keras.utils import to_categorical
from glob import glob
```

```
[2]: df = pd.read_csv('filename_to_category_map_surnamesIJKLM.csv')
```

```
[3]: df
```

```
[3]:      image_id          cell_type  is_benign localization
0    ISIC_0027419  Benign keratosis-like lesions      1.0      scalp
1    ISIC_0026769  Benign keratosis-like lesions      1.0      scalp
2    ISIC_0031633  Benign keratosis-like lesions      1.0        ear
3    ISIC_0029176  Benign keratosis-like lesions      1.0      face
4    ISIC_0029068  Benign keratosis-like lesions      1.0      face
...
5371  ISIC_0032134          Melanocytic nevi      1.0      NaN
5372  ISIC_0032861  Benign keratosis-like lesions      1.0      NaN
5373  ISIC_0026054  Benign keratosis-like lesions      1.0      NaN
5374  ISIC_0032052            Melanoma      1.0      NaN
5375  ISIC_0031933  Benign keratosis-like lesions      1.0      NaN
```

[5376 rows x 4 columns]

```
[4]: df.dtypes
```

```
[4]: image_id      object
cell_type      object
is_benign     float64
localization    object
dtype: object
```

```
[5]: df.describe()
```

```
[5]:      is_benign
count    5276.000000
mean     0.807051
std      0.394651
min     0.000000
25%     1.000000
50%     1.000000
75%     1.000000
max     1.000000
```

```
[30]: df.isnull().sum()
```

```
[30]: image_id      0
cell_type     100
is_benign     100
localization   18
dtype: int64
```

```
[31]: df1 = df.dropna()
df1.isnull().sum()
```

```
[31]: image_id      0
cell_type     0
is_benign     0
localization   0
dtype: int64
```

```
[ ]:
```

```
[32]: df1.isnull().sum()
df1.to_csv('exclusion_dataset_task2.csv')
```

```
[33]: base_skin_dir = './ham_data_surnamesIJKLM'

imageid_path_dict = {os.path.splitext(os.path.basename(x))[0]: x
                     for x in glob(os.path.join(base_skin_dir, '*.jpg'))}
```

```
[34]: df1['path'] = df1['image_id'].map(imageid_path_dict.get)
df1.head()
```

```
C:\Users\wajah\AppData\Local\Temp\ipykernel_20744\2121549243.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
df1['path'] = df1['image_id'].map(imageid_path_dict.get)

[34]:      image_id          cell_type  is_benign localization \
0  ISIC_0027419  Benign keratosis-like lesions      1.0      scalp
1  ISIC_0026769  Benign keratosis-like lesions      1.0      scalp
2  ISIC_0031633  Benign keratosis-like lesions      1.0       ear
3  ISIC_0029176  Benign keratosis-like lesions      1.0      face
4  ISIC_0029068  Benign keratosis-like lesions      1.0      face

                           path
0  ./ham_data_surnamesIJKLM\ISIC_0027419.jpg
1  ./ham_data_surnamesIJKLM\ISIC_0026769.jpg
2  ./ham_data_surnamesIJKLM\ISIC_0031633.jpg
3  ./ham_data_surnamesIJKLM\ISIC_0029176.jpg
4  ./ham_data_surnamesIJKLM\ISIC_0029068.jpg
```

```
[35]: df1['image'] = df1['path'].map(lambda x: np.asarray(Image.open(x).  
                                         resize((125,100))))
```

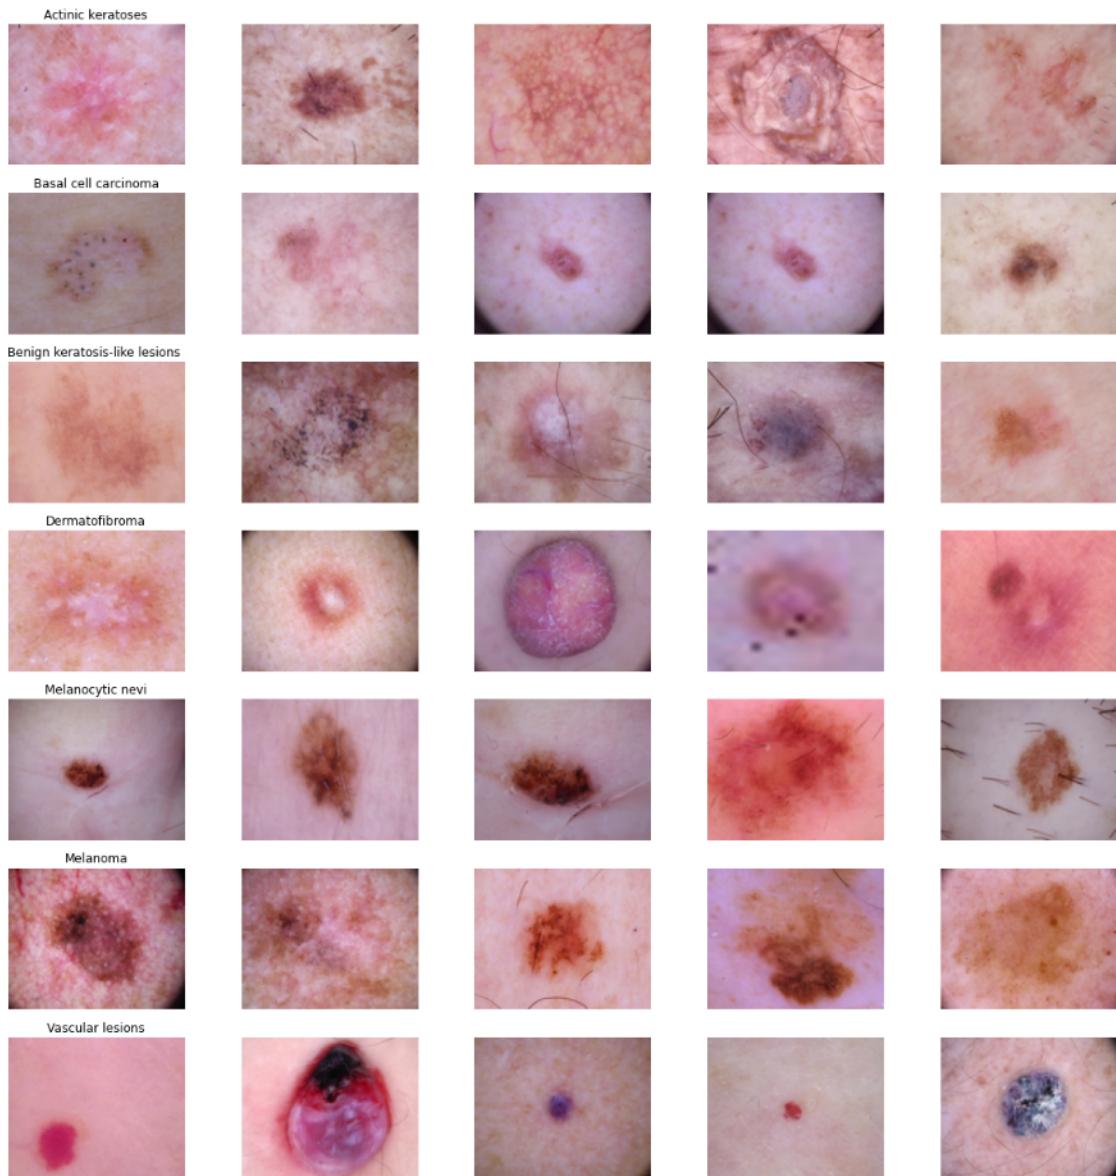
```
C:\Users\wajah\AppData\Local\Temp/ipykernel_20744/3360847188.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1['image'] = df1['path'].map(lambda x:  
np.asarray(Image.open(x).resize((125,100))))
```

```
[ ]:
```

```
[36]: n_samples = 5  
fig, m_axs = plt.subplots(7, n_samples, figsize = (4*n_samples, 3*7))  
for n_axs, (type_name, type_rows) in zip(m_axs,  
                                         df1.sort_values(['cell_type']).  
                                         groupby('cell_type')):  
    n_axs[0].set_title(type_name)  
    for c_ax, (_, c_row) in zip(n_axs, type_rows.sample(n_samples, u  
                                         random_state=2018).iterrows()):  
        c_ax.imshow(c_row['image'])  
        c_ax.axis('off')  
fig.savefig('category_samples.png', dpi=300)
```



```
[ ]: df1['image'].map(lambda x: x.shape).value_counts()
```

```
[ ]: (100, 125, 3)      5258
Name: image, dtype: int64
```

```
[ ]:
```

```
[ ]: plt.figure(figsize=(20,10))
plt.subplots_adjust(left=0.125, bottom=1, right=0.9, top=2, hspace=0.2)
```

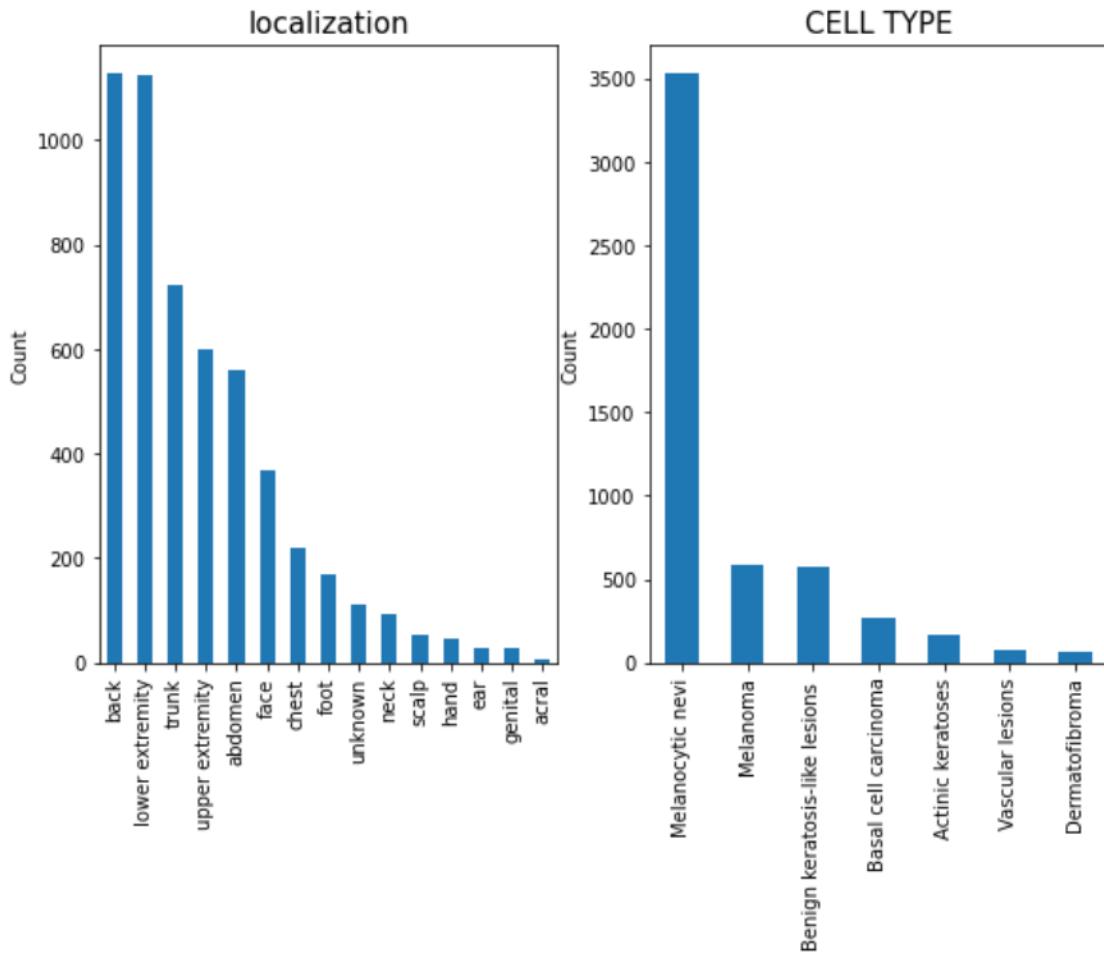
```

plt.subplot(2,4,3)
plt.title("localization", fontsize=15)
plt.ylabel("Count")
plt.xticks(rotation=45)
df1['localization'].value_counts().plot.bar()

plt.subplot(2,4,4)
plt.title("CELL TYPE", fontsize=15)
plt.ylabel("Count")
df1['cell_type'].value_counts().plot.bar()

```

[]: <AxesSubplot:title={'center':'CELL TYPE'}, ylabel='Count'>

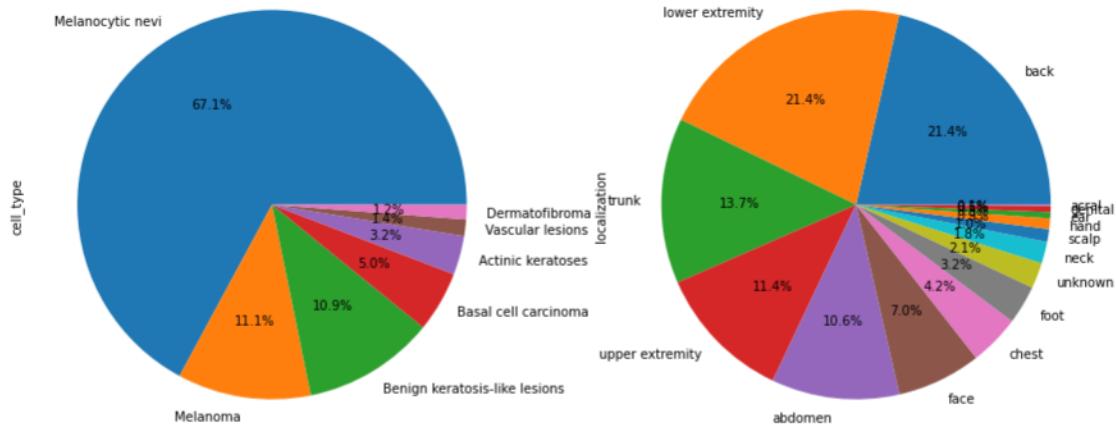


```

[ ]: plt.figure(figsize=(15,10))
plt.subplot(1,2,1)
df1['cell_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.subplot(1,2,2)

```

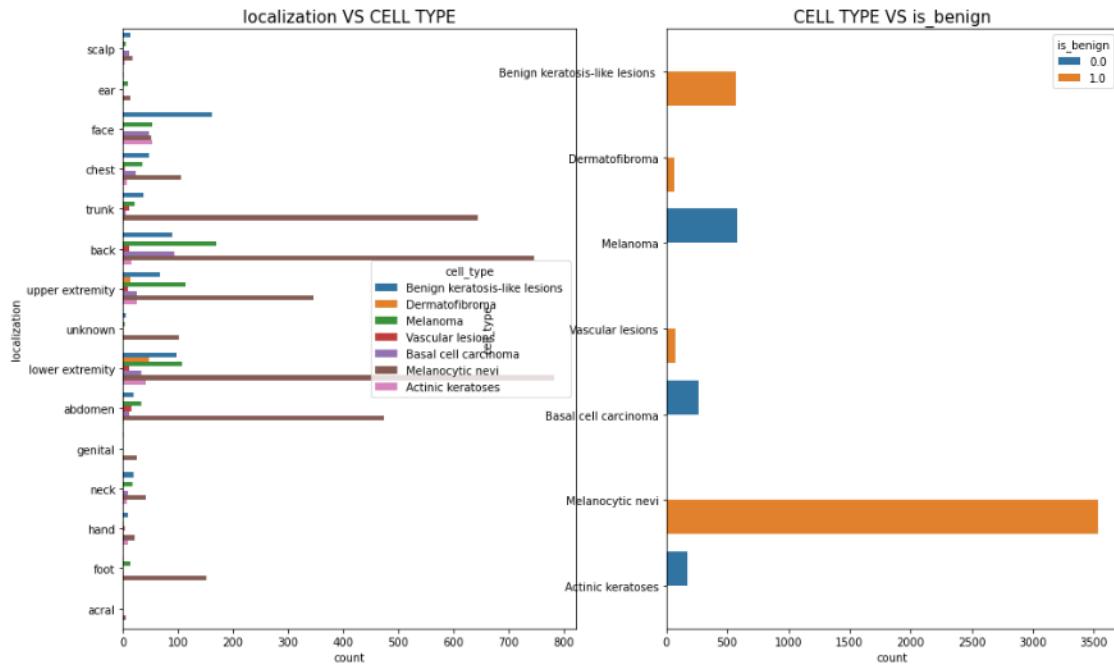
```
df1['localization'].value_counts().plot.pie(autopct="%1.1f%%")
plt.show()
```



- The face is infected the most by Benign keratosis-like lesions.
- Body parts(except face) are infected the most by Melanocytic nevi.

```
[ ]: plt.figure(figsize=(25,10))
plt.subplot(131)
plt.title('localization VS CELL TYPE', fontsize = 15)
sns.countplot(y='localization', hue='cell_type', data=df1)
plt.subplot(132)
plt.title('CELL TYPE VS is_benign', fontsize = 15)
sns.countplot(y='cell_type', hue='is_benign', data=df1)
```

```
[ ]: <AxesSubplot:title={'center':'CELL TYPE VS is_benign'}, xlabel='count',
ylabel='cell_type'>
```



```
[ ]: from sklearn.model_selection import train_test_split
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
import tensorflow as tf
from sklearn.preprocessing import StandardScaler
```

0.1 ANN

```
[ ]: features=df1.drop(columns=['is_benign'],axis=1)
target=df1['is_benign']
```

```
[ ]: features.head()
```

```
[ ]:      image_id          cell_type localization \
0  ISIC_0027419  Benign keratosis-like lesions      scalp
1  ISIC_0026769  Benign keratosis-like lesions      scalp
2  ISIC_0031633  Benign keratosis-like lesions        ear
3  ISIC_0029176  Benign keratosis-like lesions      face
4  ISIC_0029068  Benign keratosis-like lesions      face

                                         path \
0  ham_data_surnamesIJKLM\ISIC_0027419.jpg
1  ham_data_surnamesIJKLM\ISIC_0026769.jpg
2  ham_data_surnamesIJKLM\ISIC_0031633.jpg
```

```
3 ham_data_surnamesIJKLM\ISIC_0029176.jpg
4 ham_data_surnamesIJKLM\ISIC_0029068.jpg
```

```
          image
0 [[[189, 152, 194], [192, 156, 198], [191, 154,...]
1 [[[186, 127, 135], [189, 133, 145], [192, 135,...]
2 [[[131, 88, 110], [142, 97, 120], [152, 107, 1...
3 [[[190, 144, 126], [192, 145, 130], [194, 146,...]
4 [[[147, 103, 83], [153, 111, 91], [159, 119, 1...
```

```
[ ]: x_train_o, x_test_o, y_train_o, y_test_o = train_test_split(features, target,
   ↪test_size=0.25,random_state=666)
tf.unique(x_train_o.cell_type.values)
```

```
[ ]: Unique(y=<tf.Tensor: shape=(7,), dtype=string, numpy=
array([b'Melanocytic nevi', b'Basal cell carcinoma',
       b'Benign keratosis-like lesions ', b'Melanoma',
       b'Actinic keratoses', b'Dermatofibroma', b'Vascular lesions'],
      dtype=object)>, idx=<tf.Tensor: shape=(3943,), dtype=int32,
      numpy=array([0, 0, 0, ..., 0, 0, 0])>)
```

```
[ ]: x_train = np.asarray(x_train_o['image'].tolist())
x_test = np.asarray(x_test_o['image'].tolist())

x_train_mean = np.mean(x_train)
x_train_std = np.std(x_train)

x_test_mean = np.mean(x_test)
x_test_std = np.std(x_test)

x_train = (x_train - x_train_mean)/x_train_std
x_test = (x_test - x_test_mean)/x_test_std
```

```
[ ]: y_train = to_categorical(y_train_o, num_classes = 7)
y_test = to_categorical(y_test_o, num_classes = 7)
y_test
```

```
[ ]: array([[0., 1., 0., ..., 0., 0., 0.],
           [0., 1., 0., ..., 0., 0., 0.],
           [0., 1., 0., ..., 0., 0., 0.],
           ...,
           [0., 1., 0., ..., 0., 0., 0.],
           [0., 1., 0., ..., 0., 0., 0.],
           [0., 1., 0., ..., 0., 0., 0.]], dtype=float32)
```

```
[ ]: x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train,
   ↪test_size = 0.1, random_state = 999)
```

```
x_train = x_train.reshape(x_train.shape[0], *(100, 125, 3))
x_test = x_test.reshape(x_test.shape[0], *(100, 125, 3))
x_validate = x_validate.reshape(x_validate.shape[0], *(100, 125, 3))
```

```
[ ]: x_train = x_train.reshape(3548,125*100*3)
x_test = x_test.reshape(1315,125*100*3)
print(x_train.shape)
print(x_test.shape)
```

```
(3548, 37500)
(1315, 37500)
```

```
[ ]: model = Sequential()

model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu',
                ↪input_dim = 37500))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units= 64, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 7, kernel_initializer = 'uniform', activation = ↪
                ↪'softmax'))

optimizer = tf.keras.optimizers.Adam(learning_rate = 0.00075,
                                      beta_1 = 0.9,
                                      beta_2 = 0.999,
                                      epsilon = 1e-8)

model.compile(optimizer = optimizer, loss = 'categorical_crossentropy', metrics ↪
                ↪= ['accuracy'])

history = model.fit(x_train, y_train, batch_size = 10, epochs = 50)

accuracy = model.evaluate(x_test, y_test, verbose=1)[1]
print("Test: accuracy = ",accuracy*100,"%")
```

```
Epoch 1/50
355/355 [=====] - 8s 18ms/step - loss: 0.5011 -
accuracy: 0.8047
Epoch 2/50
355/355 [=====] - 8s 21ms/step - loss: 0.4329 -
accuracy: 0.8075
Epoch 3/50
355/355 [=====] - 8s 22ms/step - loss: 0.4060 -
accuracy: 0.8120 4s - loss: 0.3981 - accuracy: 0.81 - ETA: 4s - loss: 0.3
Epoch 4/50
355/355 [=====] - 9s 25ms/step - loss: 0.3898 -
```

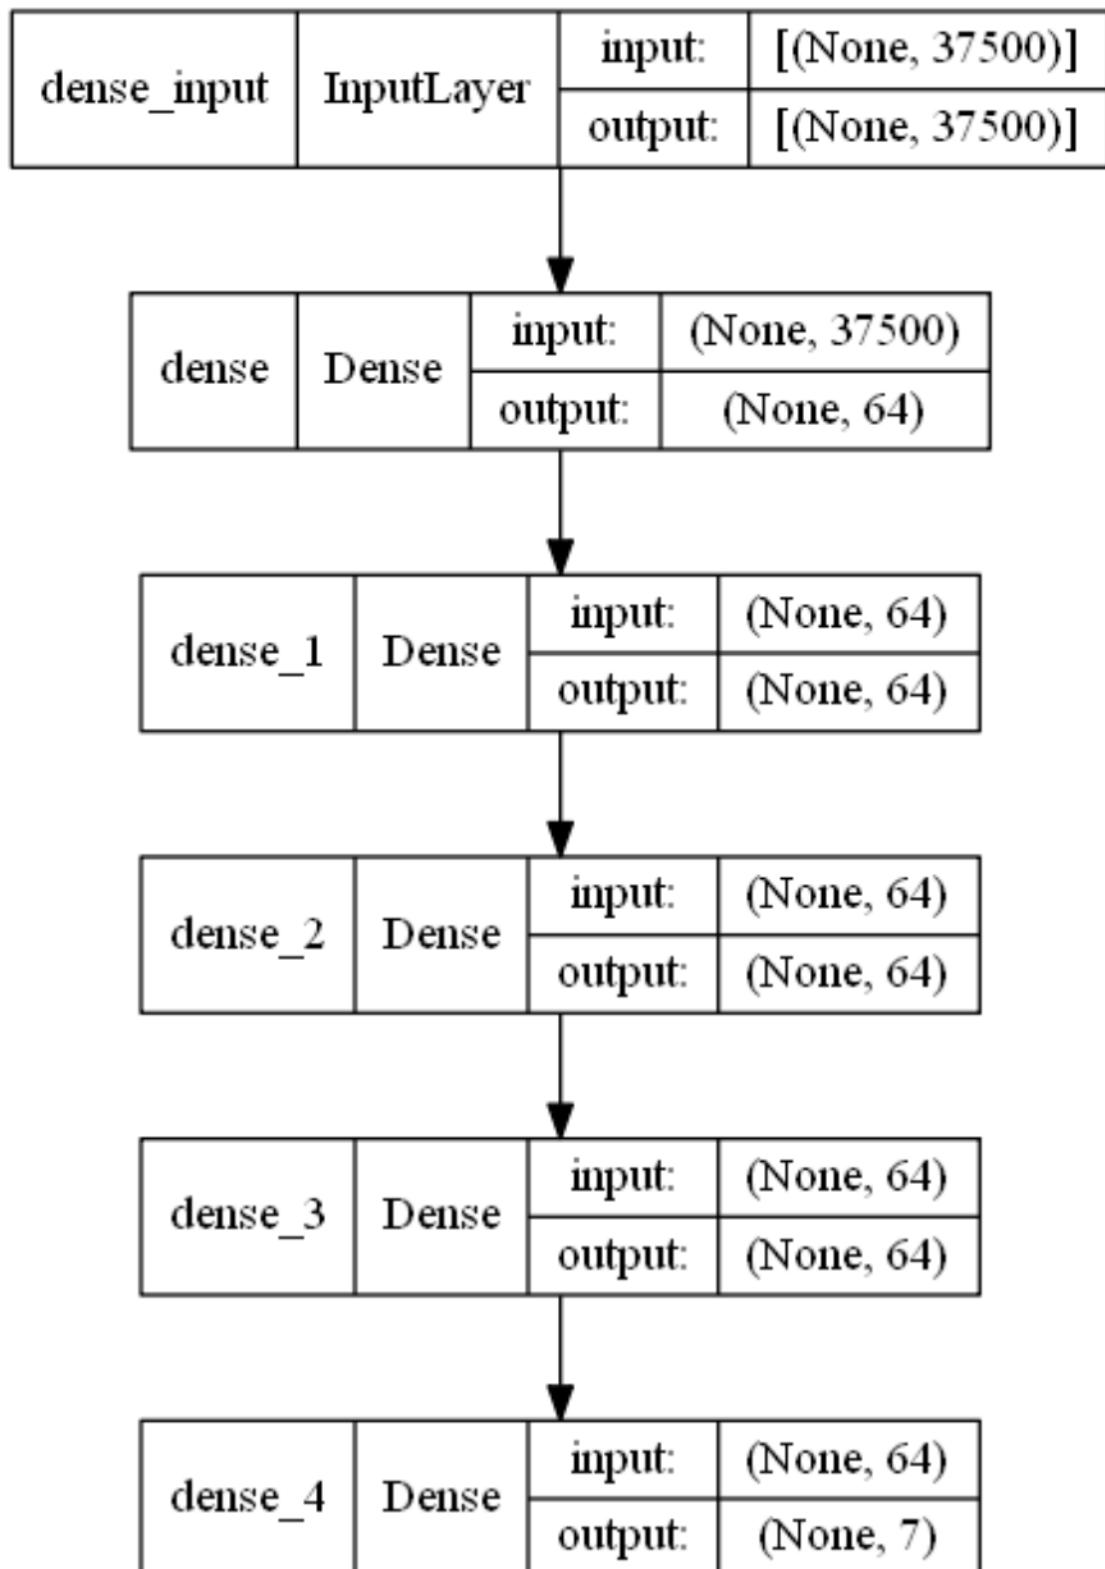
```
accuracy: 0.8247
Epoch 5/50
355/355 [=====] - 9s 24ms/step - loss: 0.3829 -
accuracy: 0.8269
Epoch 6/50
355/355 [=====] - 9s 25ms/step - loss: 0.3559 -
accuracy: 0.8391
Epoch 7/50
355/355 [=====] - 7s 20ms/step - loss: 0.3482 -
accuracy: 0.8464
Epoch 8/50
355/355 [=====] - 9s 24ms/step - loss: 0.3309 -
accuracy: 0.8591
Epoch 9/50
355/355 [=====] - 9s 25ms/step - loss: 0.3120 -
accuracy: 0.8630
Epoch 10/50
355/355 [=====] - 9s 24ms/step - loss: 0.3024 -
accuracy: 0.8701
Epoch 11/50
355/355 [=====] - 9s 25ms/step - loss: 0.2753 -
accuracy: 0.8842
Epoch 12/50
355/355 [=====] - 8s 22ms/step - loss: 0.2632 -
accuracy: 0.8943
Epoch 13/50
355/355 [=====] - 6s 17ms/step - loss: 0.2401 -
accuracy: 0.9042
Epoch 14/50
355/355 [=====] - 6s 17ms/step - loss: 0.2369 -
accuracy: 0.9030
Epoch 15/50
355/355 [=====] - 6s 17ms/step - loss: 0.2137 -
accuracy: 0.9169
Epoch 16/50
355/355 [=====] - 7s 20ms/step - loss: 0.2193 -
accuracy: 0.9101
Epoch 17/50
355/355 [=====] - 6s 18ms/step - loss: 0.1817 -
accuracy: 0.9304
Epoch 18/50
355/355 [=====] - 6s 17ms/step - loss: 0.1892 -
accuracy: 0.9278
Epoch 19/50
355/355 [=====] - 6s 18ms/step - loss: 0.1669 -
accuracy: 0.9346
Epoch 20/50
355/355 [=====] - 7s 20ms/step - loss: 0.1586 -
```

```
accuracy: 0.9360
Epoch 21/50
355/355 [=====] - 6s 16ms/step - loss: 0.1533 -
accuracy: 0.9394
Epoch 22/50
355/355 [=====] - 6s 18ms/step - loss: 0.1379 -
accuracy: 0.9490
Epoch 23/50
355/355 [=====] - 6s 17ms/step - loss: 0.1384 -
accuracy: 0.9467
Epoch 24/50
355/355 [=====] - 6s 17ms/step - loss: 0.1369 -
accuracy: 0.9467
Epoch 25/50
355/355 [=====] - 6s 17ms/step - loss: 0.1268 -
accuracy: 0.9512
Epoch 26/50
355/355 [=====] - 6s 16ms/step - loss: 0.1239 -
accuracy: 0.9498
Epoch 27/50
355/355 [=====] - 6s 17ms/step - loss: 0.1092 -
accuracy: 0.9583
Epoch 28/50
355/355 [=====] - 6s 16ms/step - loss: 0.1105 -
accuracy: 0.9583
Epoch 29/50
355/355 [=====] - 6s 17ms/step - loss: 0.0962 -
accuracy: 0.9636
Epoch 30/50
355/355 [=====] - 6s 17ms/step - loss: 0.0931 -
accuracy: 0.9673
Epoch 31/50
355/355 [=====] - 6s 18ms/step - loss: 0.0999 -
accuracy: 0.9636
Epoch 32/50
355/355 [=====] - 7s 18ms/step - loss: 0.0807 -
accuracy: 0.9710
Epoch 33/50
355/355 [=====] - 6s 17ms/step - loss: 0.0750 -
accuracy: 0.9690
Epoch 34/50
355/355 [=====] - 6s 17ms/step - loss: 0.0884 -
accuracy: 0.9676
Epoch 35/50
355/355 [=====] - 6s 17ms/step - loss: 0.0882 -
accuracy: 0.9713
Epoch 36/50
355/355 [=====] - 6s 16ms/step - loss: 0.0648 -
```

```
accuracy: 0.9741
Epoch 37/50
355/355 [=====] - 6s 16ms/step - loss: 0.0952 -
accuracy: 0.9696
Epoch 38/50
355/355 [=====] - 6s 16ms/step - loss: 0.0634 -
accuracy: 0.9766
Epoch 39/50
355/355 [=====] - 5s 15ms/step - loss: 0.0766 -
accuracy: 0.9732
Epoch 40/50
355/355 [=====] - 5s 15ms/step - loss: 0.0669 -
accuracy: 0.9760
Epoch 41/50
355/355 [=====] - 6s 16ms/step - loss: 0.0727 -
accuracy: 0.9797
Epoch 42/50
355/355 [=====] - 5s 14ms/step - loss: 0.0797 -
accuracy: 0.9760
Epoch 43/50
355/355 [=====] - 6s 16ms/step - loss: 0.0624 -
accuracy: 0.9820
Epoch 44/50
355/355 [=====] - 6s 16ms/step - loss: 0.0586 -
accuracy: 0.9794
Epoch 45/50
355/355 [=====] - 6s 16ms/step - loss: 0.0720 -
accuracy: 0.9752
Epoch 46/50
355/355 [=====] - 6s 16ms/step - loss: 0.0561 -
accuracy: 0.9825
Epoch 47/50
355/355 [=====] - 6s 16ms/step - loss: 0.0596 -
accuracy: 0.9763
Epoch 48/50
355/355 [=====] - 6s 16ms/step - loss: 0.0568 -
accuracy: 0.9822
Epoch 49/50
355/355 [=====] - 6s 16ms/step - loss: 0.0632 -
accuracy: 0.9780 0s - loss: 0.0
Epoch 50/50
355/355 [=====] - 6s 16ms/step - loss: 0.0777 -
accuracy: 0.9735
42/42 [=====] - 1s 7ms/step - loss: 1.2581 - accuracy:
0.8205
Test: accuracy = 82.05323219299316 %
```

```
[ ]: from keras.utils.vis_utils import plot_model
      from tensorflow.keras.utils import plot_model
      plot_model(model, to_file='model_plot.png', show_shapes=True,
      ↪show_layer_names=True)
```

```
[ ]:
```



0.2 CNN

```
[ ]: from tensorflow.keras.layers import  
    ↪Flatten,Dense,Dropout,BatchNormalization,Conv2D, MaxPool2D  
from tensorflow.keras.optimizers import Adam  
from keras.callbacks import ReduceLROnPlateau  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
  
[ ]: input_shape = (100, 125, 3)  
num_classes = 7  
  
model = Sequential()  
model.add(Conv2D(32, kernel_size=(3, 3),activation='relu',padding =  
    ↪'Same',input_shape=input_shape))  
model.add(Conv2D(32,kernel_size=(3, 3), activation='relu',padding = 'Same'))  
model.add(MaxPool2D(pool_size = (2, 2)))  
model.add(Dropout(0.16))  
  
model.add(Conv2D(32, kernel_size=(3, 3),activation='relu',padding = 'Same'))  
model.add(Conv2D(32,kernel_size=(3, 3), activation='relu',padding = 'Same'))  
model.add(MaxPool2D(pool_size = (2, 2)))  
model.add(Dropout(0.20))  
  
model.add(Conv2D(64, (3, 3), activation='relu',padding = 'same'))  
model.add(Conv2D(64, (3, 3), activation='relu',padding = 'Same'))  
model.add(MaxPool2D(pool_size=(2, 2)))  
model.add(Dropout(0.25))  
  
model.add(Flatten())  
model.add(Dense(256, activation='relu'))  
model.add(Dense(128, activation='relu'))  
model.add(Dropout(0.4))  
model.add(Dense(num_classes, activation='softmax'))  
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 100, 125, 32)	896
conv2d_1 (Conv2D)	(None, 100, 125, 32)	9248
max_pooling2d (MaxPooling2D	(None, 50, 62, 32)	0
)		
dropout (Dropout)	(None, 50, 62, 32)	0

conv2d_2 (Conv2D)	(None, 50, 62, 32)	9248
conv2d_3 (Conv2D)	(None, 50, 62, 32)	9248
max_pooling2d_1 (MaxPooling 2D)	(None, 25, 31, 32)	0
dropout_1 (Dropout)	(None, 25, 31, 32)	0
conv2d_4 (Conv2D)	(None, 25, 31, 64)	18496
conv2d_5 (Conv2D)	(None, 25, 31, 64)	36928
max_pooling2d_2 (MaxPooling 2D)	(None, 12, 15, 64)	0
dropout_2 (Dropout)	(None, 12, 15, 64)	0
flatten (Flatten)	(None, 11520)	0
dense_5 (Dense)	(None, 256)	2949376
dense_6 (Dense)	(None, 128)	32896
dropout_3 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 7)	903

```
=====
Total params: 3,067,239
Trainable params: 3,067,239
Non-trainable params: 0
```

[]: optimizer = Adam(lr=0.0001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgrad=False)

```
C:\Users\wajah\anaconda3\lib\site-packages\keras\optimizer_v2\adam.py:105:
UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
    super(Adam, self).__init__(name, **kwargs)
```

[]: model.compile(optimizer = optimizer , loss = "categorical_crossentropy", metrics=["accuracy"])

[]: learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
 patience=4,
 verbose=1,

```

        factor=0.5,
        min_lr=0.00001)

[ ]: x_train, x_validate, y_train, y_validate = train_test_split(x_train, y_train,
   ↪test_size = 0.1, random_state = 999)
x_train = x_train.reshape(x_train.shape[0], *(100, 125, 3))
x_test = x_test.reshape(x_test.shape[0], *(100, 125, 3))
x_validate = x_validate.reshape(x_validate.shape[0], *(100, 125, 3))

datagen = ImageDataGenerator(
    featurewise_center=False,
    samplewise_center=False,
    featurewise_std_normalization=False,
    samplewise_std_normalization=False,
    zca_whitening=False,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range=0.12,
    height_shift_range=0.12,
    horizontal_flip=True,
    vertical_flip=True)

datagen.fit(x_train)

```

```

[ ]: epochs = 60
batch_size = 16
history = model.fit_generator(datagen.flow(x_train,y_train,
   ↪batch_size=batch_size),
                               epochs = epochs, validation_data = ,
   ↪(x_validate,y_validate),
                               verbose = 1, steps_per_epoch=x_train.shape[0] // ,
   ↪batch_size
                               , callbacks=[learning_rate_reduction])

from tensorflow.keras.metrics import Recall
from sklearn.metrics import classification_report,confusion_matrix

```

C:\Users\wajah\AppData\Local\Temp\ipykernel_7072/87963123.py:4: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.

```

history = model.fit_generator(datagen.flow(x_train,y_train,
batch_size=batch_size),
Epoch 1/60
199/199 [=====] - 126s 629ms/step - loss: 0.5866 -
accuracy: 0.7919 - val_loss: 0.5652 - val_accuracy: 0.8113 - lr: 1.0000e-04
Epoch 2/60

```

```
199/199 [=====] - 113s 568ms/step - loss: 0.5008 -  
accuracy: 0.8023 - val_loss: 0.4964 - val_accuracy: 0.8113 - lr: 1.0000e-04  
Epoch 3/60  
199/199 [=====] - 113s 570ms/step - loss: 0.4727 -  
accuracy: 0.8017 - val_loss: 0.5108 - val_accuracy: 0.8113 - lr: 1.0000e-04  
Epoch 4/60  
199/199 [=====] - 113s 569ms/step - loss: 0.4502 -  
accuracy: 0.8011 - val_loss: 0.4127 - val_accuracy: 0.8113 - lr: 1.0000e-04  
Epoch 5/60  
199/199 [=====] - ETA: 0s - loss: 0.4285 - accuracy:  
0.8061  
Epoch 00005: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.  
199/199 [=====] - 112s 565ms/step - loss: 0.4285 -  
accuracy: 0.8061 - val_loss: 0.4019 - val_accuracy: 0.8113 - lr: 1.0000e-04  
Epoch 6/60  
199/199 [=====] - 113s 567ms/step - loss: 0.4099 -  
accuracy: 0.8067 - val_loss: 0.3872 - val_accuracy: 0.8085 - lr: 5.0000e-05  
Epoch 7/60  
199/199 [=====] - 113s 569ms/step - loss: 0.4052 -  
accuracy: 0.8001 - val_loss: 0.3824 - val_accuracy: 0.8197 - lr: 5.0000e-05  
Epoch 8/60  
199/199 [=====] - 113s 568ms/step - loss: 0.4018 -  
accuracy: 0.8111 - val_loss: 0.3852 - val_accuracy: 0.8085 - lr: 5.0000e-05  
Epoch 9/60  
199/199 [=====] - 113s 568ms/step - loss: 0.3900 -  
accuracy: 0.8061 - val_loss: 0.3776 - val_accuracy: 0.8197 - lr: 5.0000e-05  
Epoch 10/60  
199/199 [=====] - 113s 567ms/step - loss: 0.3901 -  
accuracy: 0.8048 - val_loss: 0.3829 - val_accuracy: 0.8085 - lr: 5.0000e-05  
Epoch 11/60  
199/199 [=====] - ETA: 0s - loss: 0.3885 - accuracy:  
0.8149  
Epoch 00011: ReduceLROnPlateau reducing learning rate to 2.499999936844688e-05.  
199/199 [=====] - 113s 569ms/step - loss: 0.3885 -  
accuracy: 0.8149 - val_loss: 0.3723 - val_accuracy: 0.8169 - lr: 5.0000e-05  
Epoch 12/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3902 -  
accuracy: 0.8130 - val_loss: 0.3709 - val_accuracy: 0.8225 - lr: 2.5000e-05  
Epoch 13/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3796 -  
accuracy: 0.8118 - val_loss: 0.3708 - val_accuracy: 0.8169 - lr: 2.5000e-05  
Epoch 14/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3820 -  
accuracy: 0.8137 - val_loss: 0.3743 - val_accuracy: 0.8141 - lr: 2.5000e-05  
Epoch 15/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3781 -  
accuracy: 0.8171 - val_loss: 0.3644 - val_accuracy: 0.8366 - lr: 2.5000e-05  
Epoch 16/60
```

```
199/199 [=====] - 112s 561ms/step - loss: 0.3768 -  
accuracy: 0.8174 - val_loss: 0.3711 - val_accuracy: 0.8141 - lr: 2.5000e-05  
Epoch 17/60  
199/199 [=====] - 112s 562ms/step - loss: 0.3777 -  
accuracy: 0.8174 - val_loss: 0.3614 - val_accuracy: 0.8338 - lr: 2.5000e-05  
Epoch 18/60  
199/199 [=====] - 112s 565ms/step - loss: 0.3723 -  
accuracy: 0.8196 - val_loss: 0.3674 - val_accuracy: 0.8310 - lr: 2.5000e-05  
Epoch 19/60  
199/199 [=====] - ETA: 0s - loss: 0.3727 - accuracy:  
0.8149  
Epoch 00019: ReduceLROnPlateau reducing learning rate to 1.249999968422344e-05.  
199/199 [=====] - 112s 565ms/step - loss: 0.3727 -  
accuracy: 0.8149 - val_loss: 0.3679 - val_accuracy: 0.8338 - lr: 2.5000e-05  
Epoch 20/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3721 -  
accuracy: 0.8155 - val_loss: 0.3587 - val_accuracy: 0.8394 - lr: 1.2500e-05  
Epoch 21/60  
199/199 [=====] - 112s 562ms/step - loss: 0.3659 -  
accuracy: 0.8140 - val_loss: 0.3718 - val_accuracy: 0.8310 - lr: 1.2500e-05  
Epoch 22/60  
199/199 [=====] - 112s 562ms/step - loss: 0.3657 -  
accuracy: 0.8171 - val_loss: 0.3630 - val_accuracy: 0.8310 - lr: 1.2500e-05  
Epoch 23/60  
199/199 [=====] - 120s 604ms/step - loss: 0.3673 -  
accuracy: 0.8162 - val_loss: 0.3622 - val_accuracy: 0.8310 - lr: 1.2500e-05  
Epoch 24/60  
199/199 [=====] - ETA: 0s - loss: 0.3641 - accuracy:  
0.8209  
Epoch 00024: ReduceLROnPlateau reducing learning rate to 1e-05.  
199/199 [=====] - 113s 565ms/step - loss: 0.3641 -  
accuracy: 0.8209 - val_loss: 0.3600 - val_accuracy: 0.8394 - lr: 1.2500e-05  
Epoch 25/60  
199/199 [=====] - 112s 561ms/step - loss: 0.3598 -  
accuracy: 0.8294 - val_loss: 0.3596 - val_accuracy: 0.8366 - lr: 1.0000e-05  
Epoch 26/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3692 -  
accuracy: 0.8196 - val_loss: 0.3582 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 27/60  
199/199 [=====] - 113s 565ms/step - loss: 0.3622 -  
accuracy: 0.8234 - val_loss: 0.3592 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 28/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3641 -  
accuracy: 0.8137 - val_loss: 0.3592 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 29/60  
199/199 [=====] - 113s 565ms/step - loss: 0.3622 -  
accuracy: 0.8234 - val_loss: 0.3700 - val_accuracy: 0.8338 - lr: 1.0000e-05  
Epoch 30/60
```

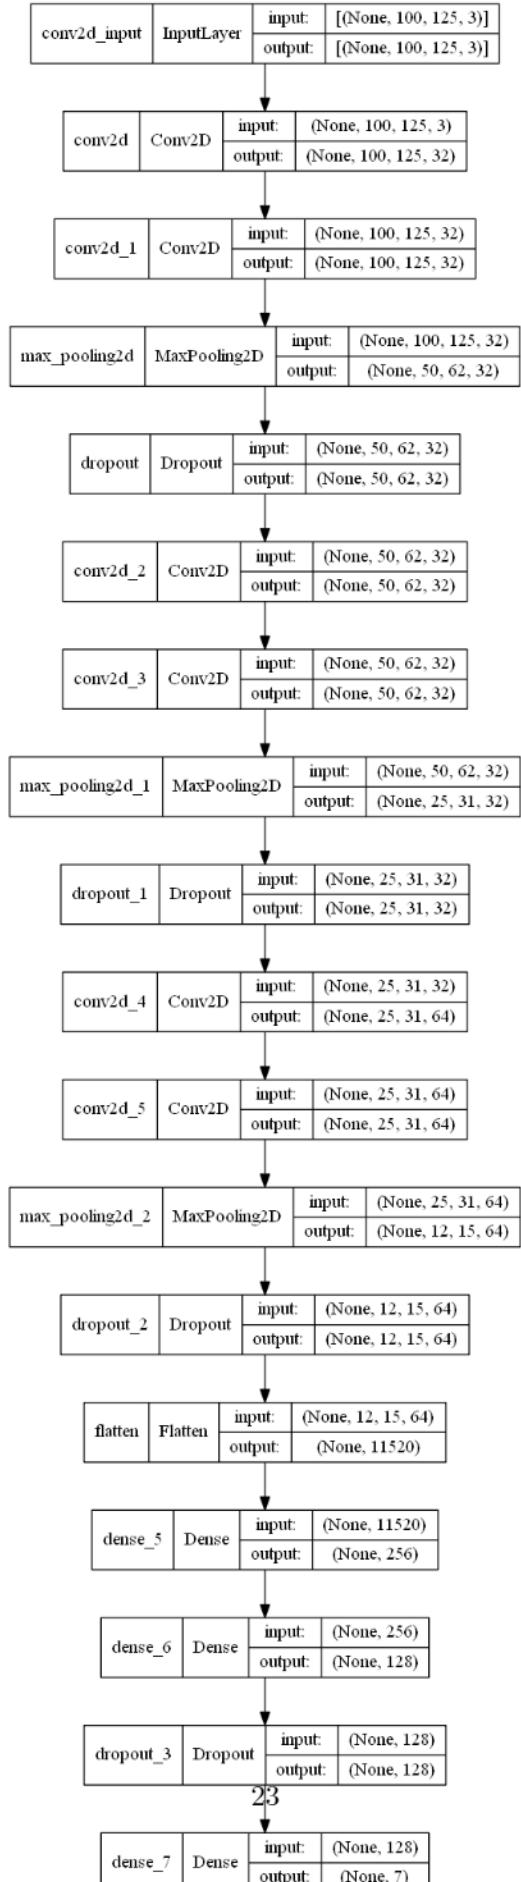
```
199/199 [=====] - 112s 562ms/step - loss: 0.3617 -  
accuracy: 0.8209 - val_loss: 0.3536 - val_accuracy: 0.8366 - lr: 1.0000e-05  
Epoch 31/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3593 -  
accuracy: 0.8181 - val_loss: 0.3534 - val_accuracy: 0.8366 - lr: 1.0000e-05  
Epoch 32/60  
199/199 [=====] - 112s 561ms/step - loss: 0.3559 -  
accuracy: 0.8228 - val_loss: 0.3563 - val_accuracy: 0.8338 - lr: 1.0000e-05  
Epoch 33/60  
199/199 [=====] - 115s 578ms/step - loss: 0.3569 -  
accuracy: 0.8234 - val_loss: 0.3570 - val_accuracy: 0.8366 - lr: 1.0000e-05  
Epoch 34/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3567 -  
accuracy: 0.8222 - val_loss: 0.3615 - val_accuracy: 0.8366 - lr: 1.0000e-05  
Epoch 35/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3562 -  
accuracy: 0.8225 - val_loss: 0.3546 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 36/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3597 -  
accuracy: 0.8190 - val_loss: 0.3619 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 37/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3575 -  
accuracy: 0.8234 - val_loss: 0.3555 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 38/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3538 -  
accuracy: 0.8291 - val_loss: 0.3603 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 39/60  
199/199 [=====] - 113s 566ms/step - loss: 0.3575 -  
accuracy: 0.8228 - val_loss: 0.3568 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 40/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3561 -  
accuracy: 0.8250 - val_loss: 0.3585 - val_accuracy: 0.8451 - lr: 1.0000e-05  
Epoch 41/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3551 -  
accuracy: 0.8215 - val_loss: 0.3522 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 42/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3551 -  
accuracy: 0.8275 - val_loss: 0.3569 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 43/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3529 -  
accuracy: 0.8247 - val_loss: 0.3566 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 44/60  
199/199 [=====] - 112s 564ms/step - loss: 0.3560 -  
accuracy: 0.8253 - val_loss: 0.3608 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 45/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3550 -  
accuracy: 0.8237 - val_loss: 0.3556 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 46/60
```

```
199/199 [=====] - 112s 563ms/step - loss: 0.3488 -  
accuracy: 0.8341 - val_loss: 0.3608 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 47/60  
199/199 [=====] - 112s 561ms/step - loss: 0.3505 -  
accuracy: 0.8332 - val_loss: 0.3534 - val_accuracy: 0.8451 - lr: 1.0000e-05  
Epoch 48/60  
199/199 [=====] - 111s 558ms/step - loss: 0.3505 -  
accuracy: 0.8250 - val_loss: 0.3497 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 49/60  
199/199 [=====] - 111s 560ms/step - loss: 0.3505 -  
accuracy: 0.8281 - val_loss: 0.3543 - val_accuracy: 0.8451 - lr: 1.0000e-05  
Epoch 50/60  
199/199 [=====] - 111s 559ms/step - loss: 0.3452 -  
accuracy: 0.8341 - val_loss: 0.3667 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 51/60  
199/199 [=====] - 111s 559ms/step - loss: 0.3517 -  
accuracy: 0.8269 - val_loss: 0.3552 - val_accuracy: 0.8507 - lr: 1.0000e-05  
Epoch 52/60  
199/199 [=====] - 111s 560ms/step - loss: 0.3481 -  
accuracy: 0.8332 - val_loss: 0.3563 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 53/60  
199/199 [=====] - 111s 559ms/step - loss: 0.3537 -  
accuracy: 0.8269 - val_loss: 0.3504 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 54/60  
199/199 [=====] - 112s 562ms/step - loss: 0.3514 -  
accuracy: 0.8291 - val_loss: 0.3519 - val_accuracy: 0.8394 - lr: 1.0000e-05  
Epoch 55/60  
199/199 [=====] - 118s 590ms/step - loss: 0.3540 -  
accuracy: 0.8275 - val_loss: 0.3544 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 56/60  
199/199 [=====] - 113s 566ms/step - loss: 0.3508 -  
accuracy: 0.8288 - val_loss: 0.3583 - val_accuracy: 0.8423 - lr: 1.0000e-05  
Epoch 57/60  
199/199 [=====] - 111s 559ms/step - loss: 0.3451 -  
accuracy: 0.8322 - val_loss: 0.3593 - val_accuracy: 0.8479 - lr: 1.0000e-05  
Epoch 58/60  
199/199 [=====] - 112s 562ms/step - loss: 0.3522 -  
accuracy: 0.8218 - val_loss: 0.3578 - val_accuracy: 0.8451 - lr: 1.0000e-05  
Epoch 59/60  
199/199 [=====] - 112s 563ms/step - loss: 0.3484 -  
accuracy: 0.8347 - val_loss: 0.3604 - val_accuracy: 0.8451 - lr: 1.0000e-05  
Epoch 60/60  
199/199 [=====] - 111s 560ms/step - loss: 0.3521 -  
accuracy: 0.8303 - val_loss: 0.3562 - val_accuracy: 0.8451 - lr: 1.0000e-05
```

```
[ ]: from keras.utils.vis_utils import plot_model
```

```
plot_model(model, to_file='model_plot.png', show_shapes=True,  
           show_layer_names=True)
```

[]:



```
[ ]: loss, accuracy = model.evaluate(x_test, y_test, verbose=1)
loss_v, accuracy_v = model.evaluate(x_validate, y_validate, verbose=1)
print("Validation: accuracy = %f ; loss_v = %f" % (accuracy_v, loss_v))
print("Test: accuracy = %f ; loss = %f" % (accuracy, loss))
model.save("model.h5")
```

```
42/42 [=====] - 10s 231ms/step - loss: 0.3513 -
accuracy: 0.8456
12/12 [=====] - 3s 222ms/step - loss: 0.3562 -
accuracy: 0.8451
Validation: accuracy = 0.845070 ; loss_v = 0.356162
Test: accuracy = 0.845627 ; loss = 0.351293
```

```
[ ]: import itertools
# Function to plot confusion matrix
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

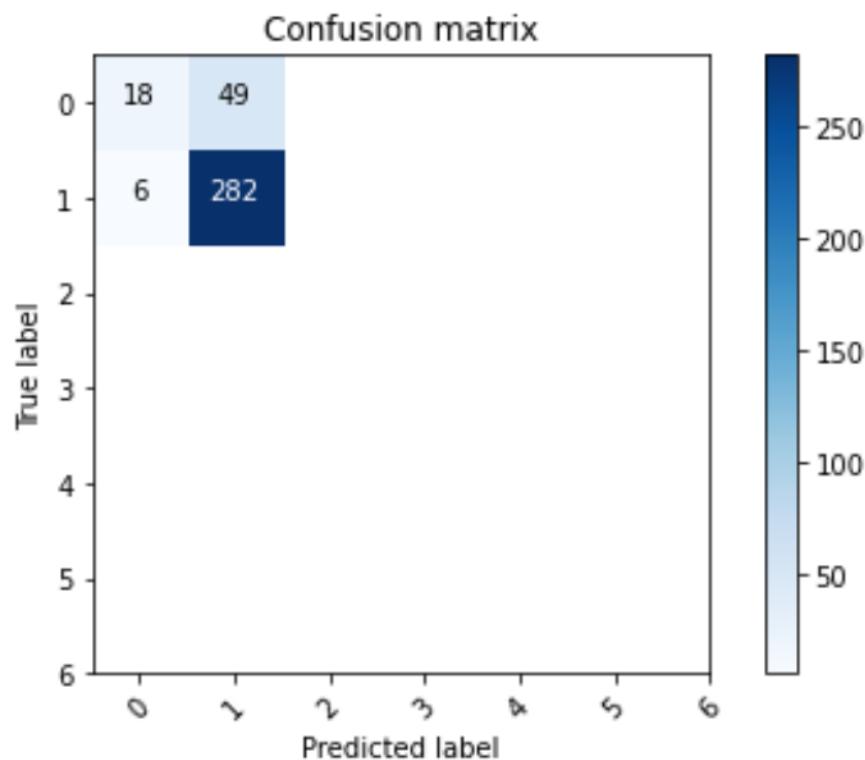
# Predict the values from the validation dataset
Y_pred = model.predict(x_validate)
```

```

# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred, axis = 1)
# Convert validation observations to one hot vectors
Y_true = np.argmax(y_validate, axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)

# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes = range(7))

```



```

[ ]: # Predict the values from the validation dataset
Y_pred = model.predict(x_test)
# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred, axis = 1)
# Convert validation observations to one hot vectors
Y_true = np.argmax(y_test, axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)

```

```
# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes = range(7))
```

